

# Experiment 10: Solving a Markov Decision Process (MDP)

[Student Name: **Anish Kumar** Student ID: **23/CS/057** Course: **BTech CSE**]

## 1. MDP and GridWorld Definition

The first task was to define the Markov Decision Process (MDP) for a 3x4 GridWorld environment. This involved setting up the core components of the model:

- **States (S):** A list of all valid (row, col) tuples in the 3x4 grid, excluding the wall at state (1, 1). Terminal states were defined as the **Goal (0, 3)** and the **Pit (1, 3)**.
- **Actions (A):** A dictionary mapping the four possible actions ('up', 'down', 'left', 'right') to their corresponding (dr, dc) movements.
- **Rewards (R):** A reward function was built, assigning **+1** to the Goal state, **-1** to the Pit state, and a **-0.04** "living penalty" to all other non-terminal states to encourage finding the shortest path.
- **Transition Model (T):** A stochastic transition model was implemented in the `get_next_states` function. For any chosen action, the agent has an **80%** chance of moving in the intended direction, a **10%** chance of "slipping" 90 degrees left, and a **10%** chance of slipping 90 degrees right. If any move results in hitting a wall or the grid boundary, the agent stays in its current state.

## 2. Core Implementation (From Scratch)

The following code snippets implement the Value Iteration and Policy Extraction algorithms as required by the lab.

## Code Snippet: Value Iteration

```
def value_iteration(states, actions, get_next_states_fn, rewards, gamma=0.99, theta=1e-4):
    V = {s: 0.0 for s in states}

    iteration = 0
    while True:
        delta = 0.0
        iteration += 1
        V_prev = V.copy()

        for s in states:
            if is_terminal(s):
                V[s] = rewards[s]
                continue

            q_values = []
            for a in actions:
                q_sa = 0.0
                for prob, s_next in get_next_states_fn(s, a):
                    r = rewards[s_next]
                    q_sa += prob * (r + gamma * V_prev[s_next])
                q_values.append(q_sa)

            best_q = max(q_values)
            V[s] = best_q
            delta = max(delta, abs(V_prev[s] - V[s]))

        if delta < theta:
            # print("Value iteration converged after {} iterations (delta={:.6f})".format(iteration, delta))
            break

    return V
```

## Code Snippet: Extract Policy

```
def extract_policy(states, actions, get_next_states_fn, rewards, V, gamma=0.99):

    policy = {}
    for s in states:
        if is_terminal(s):
            policy[s] = None
            continue

        best_a = None
        best_q = -np.inf
        for a in actions:
            q_sa = 0.0
            for prob, s_next in get_next_states_fn(s, a):
                r = rewards[s_next]
                q_sa += prob * (r + gamma * V[s_next])
            if q_sa > best_q:
                best_q = q_sa
                best_a = a
        policy[s] = best_a
    return policy
```

### 3. Experiment Results and Analysis

Three experiments were run by varying the "living penalty" (the reward for non-terminal states).

#### Experiment 1: Default Living Penalty (-0.04)

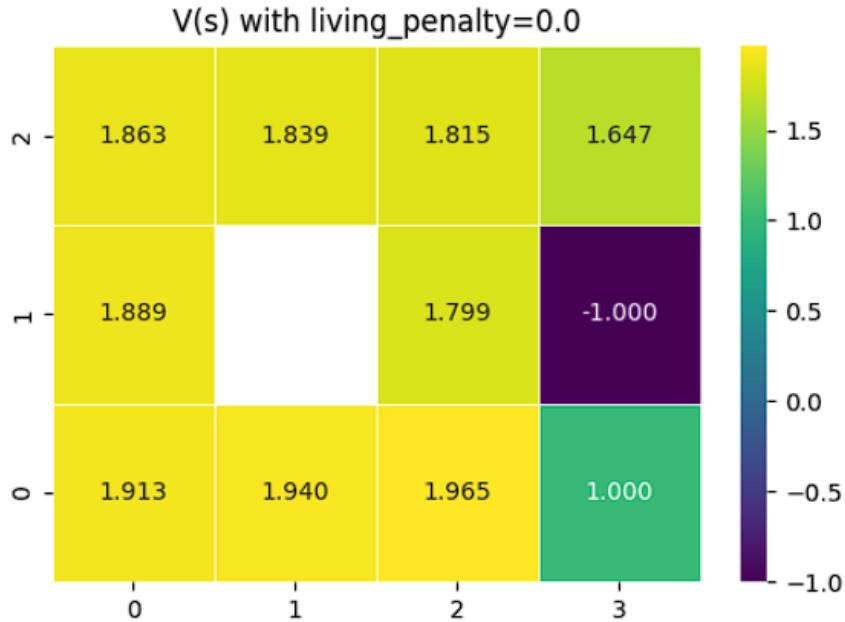
- Final Value Function (V):



(Heatmap for  $V(s)$  with living penalty=0.04)

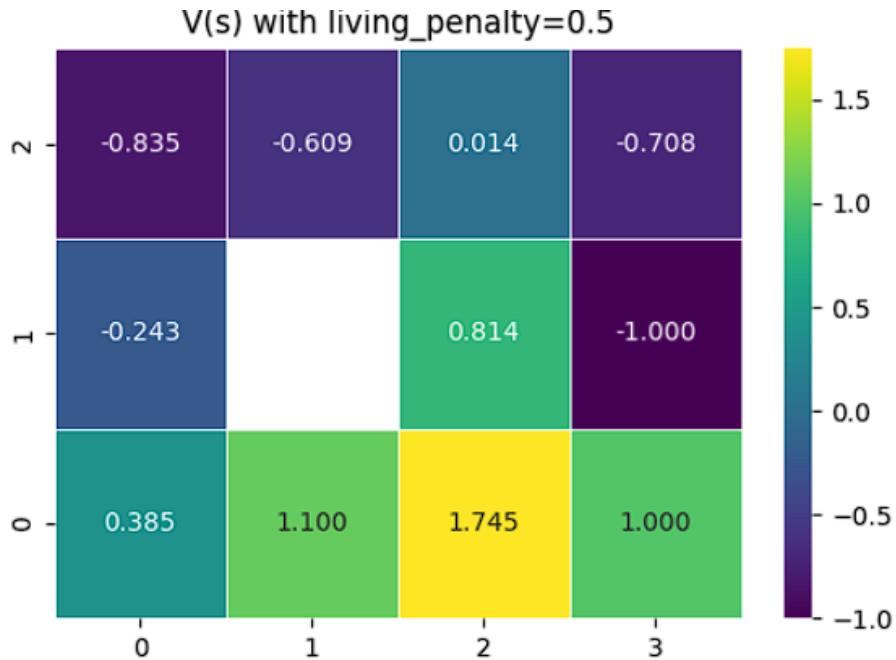
- Extracted Policy ( $\Pi_i$ ):
  - $>>> G$
  - $\wedge \# \wedge P$
  - $\wedge <<<$
- **Analysis:** The policy makes perfect sense. With a small penalty for each step, the agent learns the shortest possible path to the Goal (G) from every state. For example, from (2,0), the optimal path is  $\wedge, \wedge, >, >, >$ . The policy also correctly learns to move away from the Pit (P) and never enters the Wall (#).

#### Experiment 2: No Living Penalty (0.0)



- **Extracted Policy ( $\Pi$ ):**
- $>>> G$
- $\wedge \# < P$
- $\wedge < < V$
- **Analysis:** Yes, the policy changed. With no cost for taking steps (living\_penalty=0.0), the agent's only goal is to maximize the final reward (i.e., get +1 and avoid -1), regardless of how long it takes.
  - At state (1,2), the policy changes from  $\wedge$  (up) to  $<$  (left). The up action has a 10% chance of slipping right into the Pit. The left action (moving into the wall) has a 0% chance of reaching the pit. The agent prefers the safer, risk-free action even if it means staying in place 80% of the time.
  - At state (2,3), the policy changes from  $<$  (left) to  $v$  (down), which also results in staying in the same state, as it is safer than moving left and risking a slip into the pit.

### Experiment 3: High Living Penalty (-0.5)



- **Extracted Policy ( $\Pi$ ):**
  - $> > > G$
  - $\wedge \# \wedge P$
  - $\wedge > \wedge <$
- **Analysis:** Yes, the policy changed again. With a very high cost for each step, the agent becomes desperate to reach a terminal state (either  $G$  or  $P$ ) as quickly as possible.
  - At state  $(2,1)$ , the policy changes from  $<$  (left) to  $>$  (right). The  $>$  action moves the agent directly under the Pit. While this is closer to the negative terminal state, it is also on a shorter path (3 steps:  $\wedge, \wedge, >$ ) to the Goal than the safer 4-step path via  $(2,0)$ . The high penalty makes minimizing steps the primary objective, outweighing the slight risk.

## 4. Conclusion and Key Learnings

This lab involved implementing the Value Iteration algorithm to solve a stochastic MDP from scratch. The process provided several key insights:

1. **MDP Definition:** The lab demonstrated how to formally define an environment's dynamics. This included defining **States** (the 3x4 grid),

**Actions** (up, down, left, right), a **Reward Function** ( $R(s)$ ), and a stochastic **Transition Model** ( $T(s, a, s')$ ) with 80/10/10 probabilities.

2. **Value Iteration Implementation:** The core of the lab was implementing the Bellman Optimality Equation. The `value_iteration` function successfully found the optimal value  $V(s)$  for every state by repeatedly updating its value based on the maximum expected future reward of its neighbors.
3. **Policy Extraction:** The lab showed that once the optimal value function  $V$  is found, the optimal policy  $P_i(s)$  can be extracted. This was done with the `extract_policy` function, which performs a one-step lookahead to find the action that maximizes the expected value.
4. **Reward Function is Critical:** The experiments clearly showed that the reward function directly shapes the agent's behavior. The "living penalty" is a powerful tool to control the agent's trade-off between speed and safety.
  - **Small Penalty (-0.04):** Creates a "standard" policy that balances finding the goal quickly and avoiding the pit.
  - **No Penalty (0.0):** Makes the agent highly risk-averse. It doesn't care about path length, only about avoiding the -1 reward from the pit.
  - **High Penalty (-0.5):** Makes the agent prioritize speed above all else, taking the shortest path to a terminal state even if it's riskier (e.g., moving adjacent to the pit).